

Q1 Architecture Design

1 分數

Q1.1 Student Model Architecture

0.5 分數

Please copy&paste your student model architecture code block below. Note that the function `get_student_model()` must be included.

For example:

```
class StudentNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.cnn = nn.Sequential(
            nn.Conv2d(3, 32, 3),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.Conv2d(32, 32, 3),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(2, 2, 0),

            nn.Conv2d(32, 64, 3),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(2, 2, 0),

            nn.Conv2d(64, 100, 3),
            nn.BatchNorm2d(100),
            nn.ReLU(),
            nn.MaxPool2d(2, 2, 0),
            # Here we adopt Global Average Pooling for various input size.
            nn.AdaptiveAvgPool2d((1, 1)),
        )
        self.fc = nn.Sequential(
            nn.Linear(100, 11),
        )

    def forward(self, x):
        out = self.cnn(x)
        out = out.view(out.size()[0], -1)
        return self.fc(out)

def get_student_model():
    return StudentNet()
```

Your student model architecture **code**(text only):

```

class BasicBlock(nn.Module):
    expansion = 1
    def __init__(self, in_planes, planes, stride=1):
        super(BasicBlock, self).__init__()
        self.conv1 = deepwise_pointwise_conv(in_planes, planes, kernel_size=3,
stride=stride, padding=1)
        self.bn1 = nn.BatchNorm2d(planes)
        self.conv2 = deepwise_pointwise_conv(planes, planes, kernel_size=3,
stride=1, padding=1)
        self.bn2 = nn.BatchNorm2d(planes)

        self.shortcut = nn.Sequential()
        if stride != 1 or in_planes != self.expansion*planes:
            self.shortcut = nn.Sequential(
                deepwise_pointwise_conv(in_planes, self.expansion*planes,
kernel_size=1, stride=stride),
                nn.BatchNorm2d(self.expansion*planes)
            )

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.bn2(self.conv2(out))
        out += self.shortcut(x)
        out = F.relu(out)
        return out

```

```

class Bottleneck(nn.Module):
    expansion = 4
    def __init__(self, in_planes, planes, stride=1):
        super(Bottleneck, self).__init__()
        self.conv1 = deepwise_pointwise_conv(in_planes, planes, kernel_size=1)
        self.bn1 = nn.BatchNorm2d(planes)
        self.conv2 = deepwise_pointwise_conv(planes, planes, kernel_size=3,
stride=stride, padding=1)
        self.bn2 = nn.BatchNorm2d(planes)
        self.conv3 = deepwise_pointwise_conv(planes, self.expansion*planes,
kernel_size=1)
        self.bn3 = nn.BatchNorm2d(self.expansion*planes)

        self.shortcut = nn.Sequential()
        if stride != 1 or in_planes != self.expansion*planes:
            self.shortcut = nn.Sequential(
                deepwise_pointwise_conv(in_planes, self.expansion*planes,
kernel_size=1, stride=stride),
                nn.BatchNorm2d(self.expansion*planes)
            )

```

```

def forward(self, x):
    out = F.relu(self.bn1(self.conv1(x)))
    out = F.relu(self.bn2(self.conv2(out)))
    out = self.bn3(self.conv3(out))
    out += self.shortcut(x)
    out = F.relu(out)
    return out

class StudentNet(nn.Module):
    def __init__(self, block, num_blocks, num_classes=11):
        super().__init__()

        self.in_planes = 64
        self.conv1 = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1,
bias=False)
        self.bn1 = nn.BatchNorm2d(64)
        self.layer1 = self._make_layer(block, 64, num_blocks[0], stride=1)
        self.layer2 = self._make_layer(block, 128, num_blocks[1], stride=2)
        self.linear = nn.Linear(128*block.expansion, num_classes)

    def _make_layer(self, block, planes, num_blocks, stride):
        strides = [stride] + [1]*(num_blocks-1)
        layers = []
        for stride in strides:
            layers.append(block(self.in_planes, planes, stride))
            self.in_planes = planes * block.expansion
        return nn.Sequential(*layers)

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.layer1(out)
        out = self.layer2(out)
        out = F.adaptive_avg_pool2d(F.avg_pool2d(out, 4), 1)
        out = out.view(out.size(0), -1)
        out = self.linear(out)
        return out

    def get_student_model():
        return StudentNet(BasicBlock, [2, 2])

```

Q1.2 Torchsummary of Student Model
0.5 分數

Copy&Paste the torchsummary result of your student model. The total params should not exceed 100,000.

For example:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 222, 222]	896
BatchNorm2d-2	[-1, 32, 222, 222]	64
ReLU-3	[-1, 32, 222, 222]	0
Conv2d-4	[-1, 32, 220, 220]	9,248
BatchNorm2d-5	[-1, 32, 220, 220]	64
ReLU-6	[-1, 32, 220, 220]	0
MaxPool2d-7	[-1, 32, 110, 110]	0
Conv2d-8	[-1, 64, 108, 108]	18,496
BatchNorm2d-9	[-1, 64, 108, 108]	128
ReLU-10	[-1, 64, 108, 108]	0
MaxPool2d-11	[-1, 64, 54, 54]	0
Conv2d-12	[-1, 100, 52, 52]	57,700
BatchNorm2d-13	[-1, 100, 52, 52]	200
ReLU-14	[-1, 100, 52, 52]	0
MaxPool2d-15	[-1, 100, 26, 26]	0
AdaptiveAvgPool2d-16	[-1, 100, 1, 1]	0
Linear-17	[-1, 11]	1,111

Total params: 87,907

Trainable params: 87,907

Non-trainable params: 0

Input size (MB): 0.57

Forward/backward pass size (MB): 99.72

Params size (MB): 0.34

Estimated Total Size (MB): 100.62

Torchsummary of your student network(text only).

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 224, 224]	1,728
BatchNorm2d-2	[-1, 64, 224, 224]	128
Conv2d-3	[-1, 64, 224, 224]	640
BatchNorm2d-4	[-1, 64, 224, 224]	128
ReLU-5	[-1, 64, 224, 224]	0
Conv2d-6	[-1, 64, 224, 224]	4,160
BatchNorm2d-7	[-1, 64, 224, 224]	128
Conv2d-8	[-1, 64, 224, 224]	640
BatchNorm2d-9	[-1, 64, 224, 224]	128
ReLU-10	[-1, 64, 224, 224]	0
Conv2d-11	[-1, 64, 224, 224]	4,160

BatchNorm2d-12	[-1, 64, 224, 224]	128
BasicBlock-13	[-1, 64, 224, 224]	0
Conv2d-14	[-1, 64, 224, 224]	640
BatchNorm2d-15	[-1, 64, 224, 224]	128
ReLU-16	[-1, 64, 224, 224]	0
Conv2d-17	[-1, 64, 224, 224]	4,160
BatchNorm2d-18	[-1, 64, 224, 224]	128
Conv2d-19	[-1, 64, 224, 224]	640
BatchNorm2d-20	[-1, 64, 224, 224]	128
ReLU-21	[-1, 64, 224, 224]	0
Conv2d-22	[-1, 64, 224, 224]	4,160
BatchNorm2d-23	[-1, 64, 224, 224]	128
BasicBlock-24	[-1, 64, 224, 224]	0
Conv2d-25	[-1, 64, 112, 112]	640
BatchNorm2d-26	[-1, 64, 112, 112]	128
ReLU-27	[-1, 64, 112, 112]	0
Conv2d-28	[-1, 128, 112, 112]	8,320
BatchNorm2d-29	[-1, 128, 112, 112]	256
Conv2d-30	[-1, 128, 112, 112]	1,280
BatchNorm2d-31	[-1, 128, 112, 112]	256
ReLU-32	[-1, 128, 112, 112]	0
Conv2d-33	[-1, 128, 112, 112]	16,512
BatchNorm2d-34	[-1, 128, 112, 112]	256
Conv2d-35	[-1, 64, 112, 112]	128
BatchNorm2d-36	[-1, 64, 112, 112]	128
ReLU-37	[-1, 64, 112, 112]	0
Conv2d-38	[-1, 128, 112, 112]	8,320
BatchNorm2d-39	[-1, 128, 112, 112]	256
BasicBlock-40	[-1, 128, 112, 112]	0
Conv2d-41	[-1, 128, 112, 112]	1,280
BatchNorm2d-42	[-1, 128, 112, 112]	256
ReLU-43	[-1, 128, 112, 112]	0
Conv2d-44	[-1, 128, 112, 112]	16,512
BatchNorm2d-45	[-1, 128, 112, 112]	256
Conv2d-46	[-1, 128, 112, 112]	1,280
BatchNorm2d-47	[-1, 128, 112, 112]	256
ReLU-48	[-1, 128, 112, 112]	0
Conv2d-49	[-1, 128, 112, 112]	16,512
BatchNorm2d-50	[-1, 128, 112, 112]	256
BasicBlock-51	[-1, 128, 112, 112]	0
Linear-52	[-1, 11]	1,419
=====		
=====		
Total params: 96,587		
Trainable params: 96,587		
Non-trainable params: 0		

Input size (MB): 0.57
Forward/backward pass size (MB): 882.00
Params size (MB): 0.37
Estimated Total Size (MB): 882.94

Q2 Knowledge Distillation 1 分數

Q2.1 Knowledge Distillation with KL Divergence Loss 0.5 分數

Copy&Paste the contents in the whole code block of your `loss_fn_kd` implementation. (Simply choose `alpha=0.5`, `temperature=1.0` for the keyword arguments.)

For example:

```
def loss_fn_kd(student_logits, labels, teacher_logits, alpha=0.5, temperature=1.0):  
    # You need to implement this loss function by your own.  
    pass
```

Your `loss_fn_kd` **code**(text only).

```
def loss_fn_kd(student_logits, teacher_logits, y_a, y_b=None, lam=0, alpha=0.3,  
temperature=2.0):  
    p = F.softmax(student_logits / temperature, dim=-1)  
    q = F.softmax(teacher_logits / temperature, dim=-1)  
    kl_loss = nn.KLDivLoss(reduction='batchmean')(p, q)  
    loss_fn = nn.CrossEntropyLoss()  
    if cfg['MIXUP'] and y_b is not None:  
        ce_loss = mixup_criterion(loss_fn, student_logits, y_a, y_b, lam)  
    else:  
        ce_loss = loss_fn(student_logits, y_a)  
    loss = (alpha*temperature**2) * kl_loss + (1 - alpha) * ce_loss  
    return loss
```

//

Q2.2 Understanding Temperature in Knowledge Distillation 0.5 分數

Which is true about the hyperparameter `T` (temperature) in the knowledge distillation loss function with KL divergence loss?

Using a higher value for T produces a harder probability distribution over classes.

- ✓ Using a higher value for T produces a softer probability distribution over classes.

Q3 Network Pruning

2 分數


Q3.1 Pruning Ratio v.s. Model Accuracy

1 分數

Please go through the official pytorch network pruning tutorial first.

In this question, you are asked to plot a graph to indicate the relationship between **pruning ratio** and **model accuracy**. You can use the provided teacher model or your own teacher model. And to be more specific, please prune the parameters named with weight in all of the nn.Conv2d layers, with different ratios. If you are still confused about how to implement it, we also have provided an example on google colab.


Example plot:



Please upload your .jpg/.png file below.

▼ Q3-1.png

Download



We encourage you to do the same thing on your student network and compare the results, but there is no need to submit the plot of the student network.

Q3.2 Network Pruning and Inference Speed

1 分數

According to your understanding of the network pruning methods in the tutorial (`torch.nn.utils.prune`), can this kind of approach reduce the model inference time? Choose your answer below.

If you have no idea, we encourage you to do some experiments about it. (e.g., measure the inference time of a non-pruned model and a 90%-pruned model.)

- ☐ Yes. After pruning, the neurons are removed so the amount of calculation is lower and the inference time is faster.
- ☒ No. Such pruning technique in the tutorial is just to add some masks to the modules. The amount of calculation is nearly the same so the inference time is also similar.

HW13

● 已批改

總分

3.75 / 4 pts

問題 1

Architecture Design		1 / 1 pt
1.1	Student Model Architecture	0.5 / 0.5 pts
1.2	Torchsummary of Student Model	0.5 / 0.5 pts
問題 2		
Knowledge Distillation		0.75 / 1 pt
2.1	Knowledge Distillation with KL Divergence Loss	0.25 / 0.5 pts
2.2	Understanding Temperature in Knowledge Distillation	0.5 / 0.5 pts
問題 3		
Network Pruning		2 / 2 pts
3.1	Pruning Ratio v.s. Model Accuracy	1 / 1 pt
3.2	Network Pruning and Inference Speed	1 / 1 pt